

Lab 1: Data Visualization, Data Preprocessing, and Statistical Analysis

Nitish Dhinaharan

University of the Cumberlands

MSCS-634-B01 Advanced Big Data and Data Mining

Dr. Satish Penmatsa

November 02, 2025

Lab 1: Data Visualization, Data Preprocessing, and Statistical Analysis

Github Link

https://github.com/ndhinaharan36295/MSCS-634_Lab-1

1. Data Collection

```
Lab 1: Data Visualization, Data Preprocessing, and Statistical Analysis Using Python in Jupyter Notebook

[14]
✓ Os
▶ # Imports
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import MinMaxScaler

[4]
✓ Os
# Step 1: Data Collection

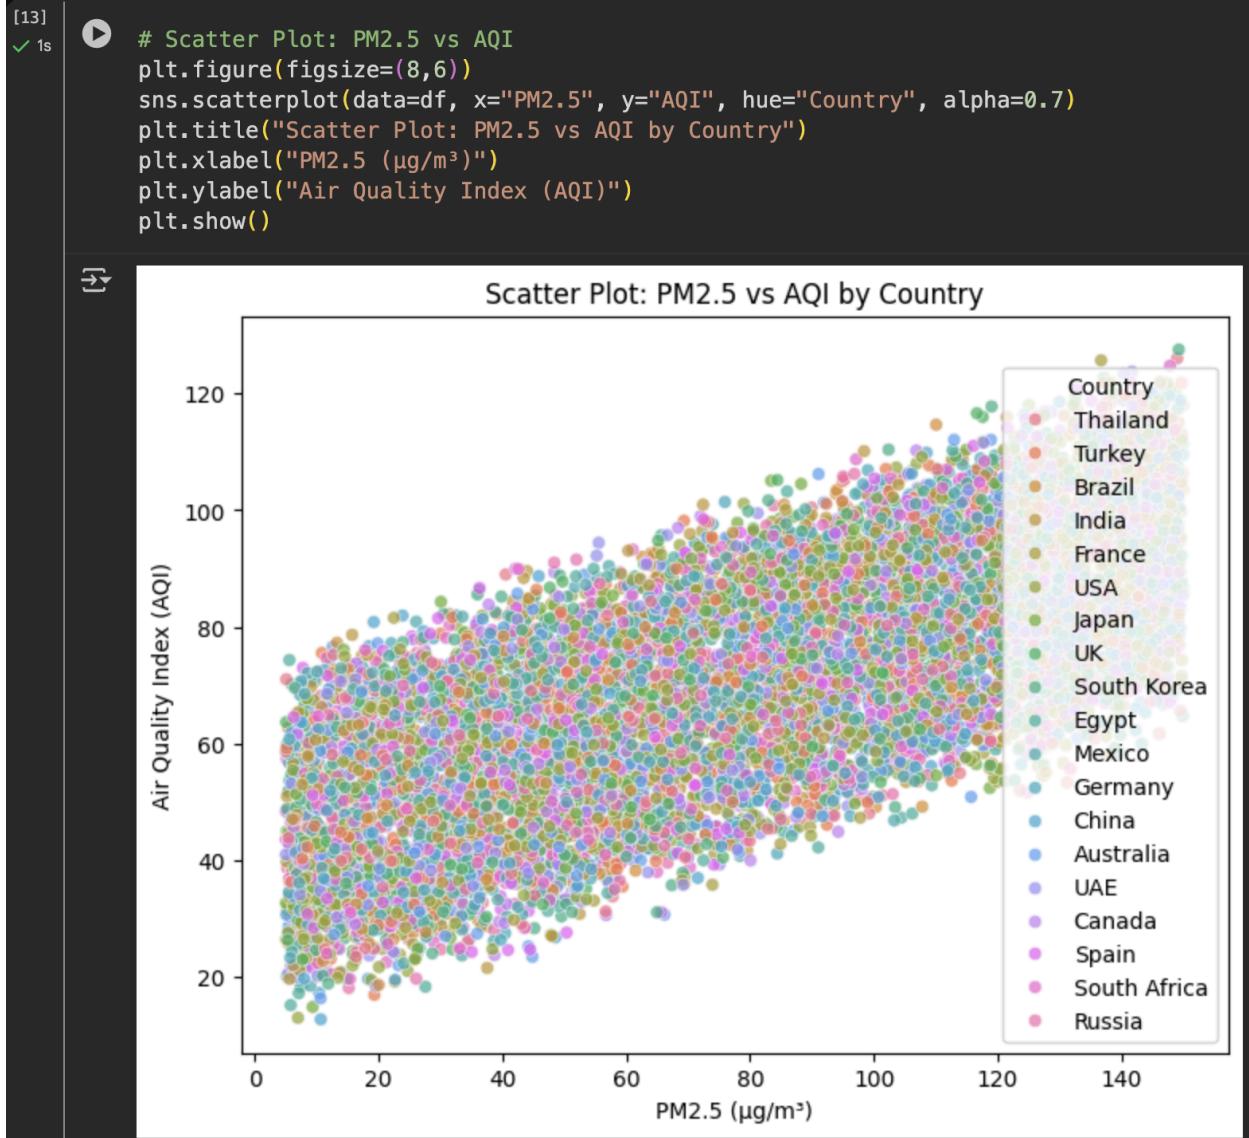
# Load dataset
df = pd.read_csv("global_air_quality.csv")

# Display first five rows
df.head()

      City  Country       Date  PM2.5  PM10   NO2   SO2    CO    O3 Temperature  Humidity  Wind Speed
0    Bangkok   Thailand 2023-03-19  86.57  25.19  99.88  30.63  4.46  36.29      17.67     59.35    13.76
1    Istanbul    Turkey 2023-02-16  50.63  97.39  48.14   8.71  3.40 144.16      3.46     67.51     6.36
2  Rio de Janeiro   Brazil 2023-11-13 130.21  57.22  98.51   9.92  0.12 179.31      25.29     29.30    12.87
3     Mumbai     India 2023-03-16 119.70 130.52  10.96  33.03  7.74  38.65      23.15     99.97     7.71
4      Paris     France 2023-04-04  55.20  36.62  76.85  21.85  2.00  67.09      16.02     90.28    14.16
```

2. Data Visualization

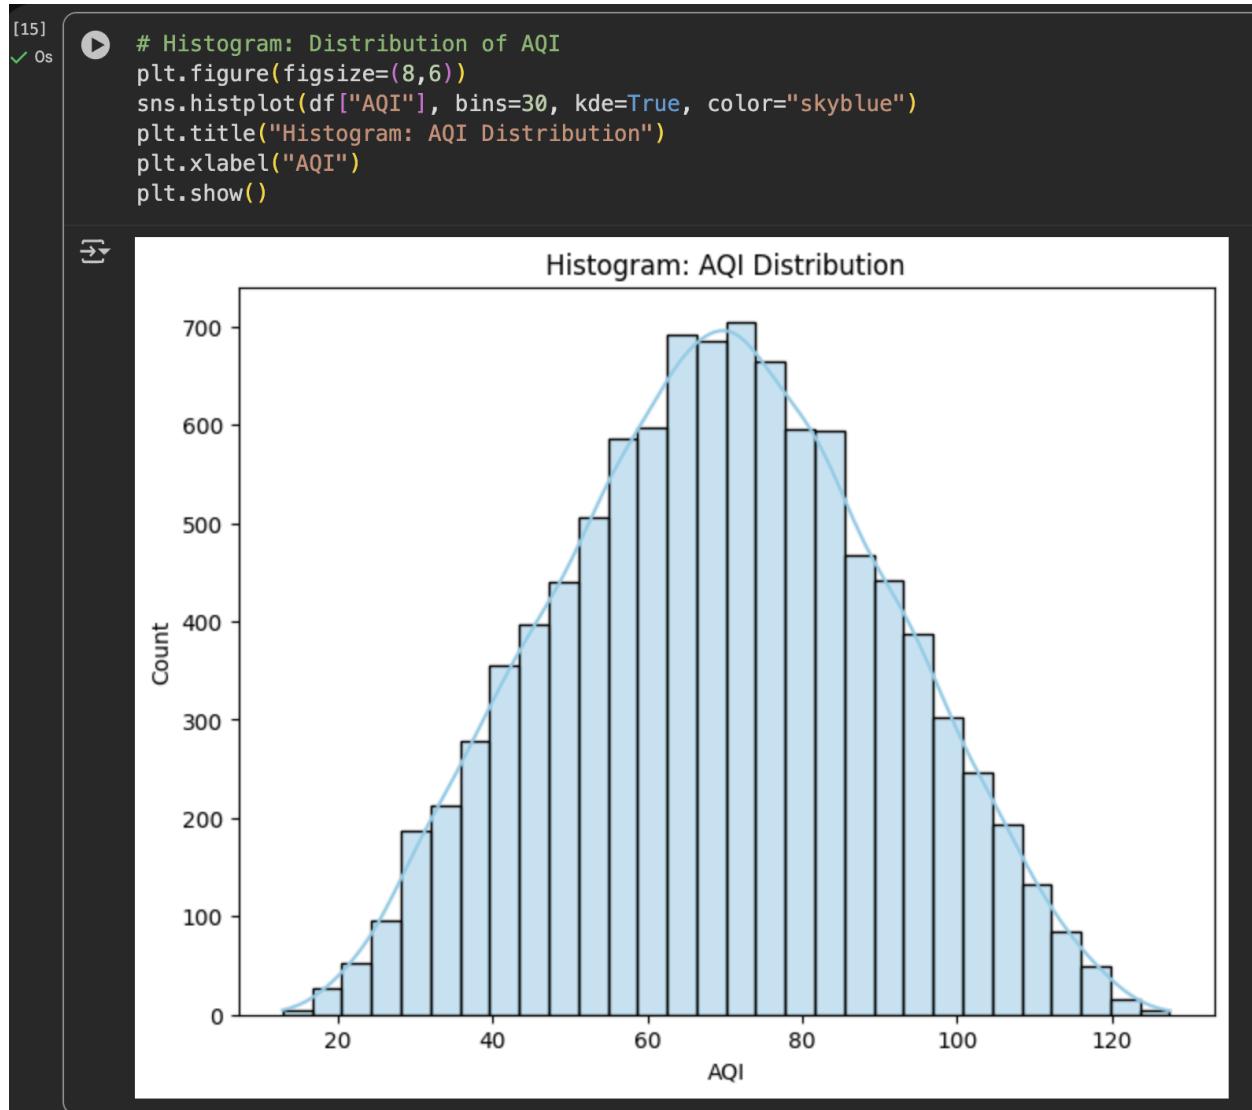
1) Scatter Plot: PM2.5 vs AQI



The scatter plot of PM2.5 versus AQI by country reveals a strong positive correlation, indicating that as PM2.5 concentrations increase, air quality deteriorates correspondingly. Most data points cluster between 20–80 $\mu\text{g}/\text{m}^3$ of PM2.5 and 40–100 AQI, representing moderate pollution levels common across regions. The consistent upward trend across countries suggests that PM2.5 is a

dominant factor influencing overall air quality worldwide. Some vertical dispersion at similar PM_{2.5} levels implies the impact of secondary pollutants or meteorological variations, while a few outliers at higher values likely reflect severe pollution episodes in specific urban or industrial areas.

2) Distribution of AQI Histogram



The histogram of AQI values shows a roughly bell-shaped distribution, indicating that most regions experience moderate air quality conditions centered around AQI values of 60 to 80. The symmetry of the distribution suggests that extreme pollution or exceptionally clean air events are relatively uncommon. The smooth, near-normal pattern also implies a stable environmental range across locations, with only a small number of observations falling into high-AQI categories that represent more severe pollution levels.

3. Data Preprocessing

1) Handling missing values

Before:

```
Step 3: Data Preprocessing

[21]
✓ Os
# Handling missing values

# Check for missing values before
print(df.isnull().sum())

City          0
Country       0
Date          0
PM2.5         0
PM10          0
NO2           0
SO2           0
CO            0
O3            0
Temperature   0
Humidity      0
Wind Speed    0
AQI           0
dtype: int64
```

After:

```
[22]  ✓ 0s
▶ # Fill missing values of pollutant & AQI columns with median
cols = ["PM2.5", "PM10", "NO2", "SO2", "CO", "O3", "AQI"]
for c in cols:
    df[c] = df[c].fillna(df[c].median())

# verify no missing values after
print(df.isnull().sum())

⇒ City          0
Country        0
Date           0
PM2.5          0
PM10           0
NO2            0
SO2            0
CO             0
O3             0
Temperature    0
Humidity       0
Wind Speed     0
AQI            0
dtype: int64
```

2) Outlier Detection & Removal (IQR method)

```
[27]  ✓ 0s
▶ # 2) Outlier Detection and Removal

# Example for PM2.5
Q1 = df["PM2.5"].quantile(0.25)
Q3 = df["PM2.5"].quantile(0.75)
IQR = Q3 - Q1
lower = Q1 - 1.5 * IQR
upper = Q3 + 1.5 * IQR

print("IQR:", IQR)

print("Before removing outliers:", df.shape)

outliers = df[(df["PM2.5"] < lower) | (df["PM2.5"] > upper)]
print("Outliers count:", outliers.shape[0])

df_no_outliers = df[(df["PM2.5"] >= lower) & (df["PM2.5"] <= upper)]
print("After removing outliers:", df_no_outliers.shape)

⇒ IQR: 72.2075
Before removing outliers: (10000, 13)
Outliers count: 0
After removing outliers: (10000, 13)
```

3) Data Reduction

[30] ✓ 0s # 3) Data Reduction
df.head()

	City	Country	Date	PM2.5	PM10	NO2	SO2	CO	O3	Temperature	Humidity	Wind Speed	AQI
0	Bangkok	Thailand	2023-03-19	86.57	25.19	99.88	30.63	4.46	36.29	17.67	59.35	13.76	56.9025
1	Istanbul	Turkey	2023-02-16	50.63	97.39	48.14	8.71	3.40	144.16	3.46	67.51	6.36	57.7080
2	Rio de Janeiro	Brazil	2023-11-13	130.21	57.22	98.51	9.92	0.12	179.31	25.29	29.30	12.87	84.6245
3	Mumbai	India	2023-03-16	119.70	130.52	10.96	33.03	7.74	38.65	23.15	99.97	7.71	82.1785
4	Paris	France	2023-04-04	55.20	36.62	76.85	21.85	2.00	67.09	16.02	90.28	14.16	45.7420

Next steps: [Generate code with df](#) [New interactive sheet](#)

[31] ✓ 0s ⏪ # Sampling (e.g., 20%)
df_sample = df_no_outliers.sample(n=int(0.2 * len(df_no_outliers)), random_state=42)

Dropping columns less relevant (e.g., "Date")
df_reduced = df_sample.drop(["Date"], axis=1) # adjust column list as needed
df_reduced.head()

	City	Country	PM2.5	PM10	NO2	SO2	CO	O3	Temperature	Humidity	Wind Speed	AQI
6252	Beijing	China	111.92	32.08	7.95	37.45	9.43	45.77	23.68	30.31	5.91	55.3610
4684	Berlin	Germany	62.62	153.88	22.76	22.12	0.45	112.53	1.93	66.61	1.22	71.6845
1731	Bangkok	Thailand	77.06	114.31	24.99	45.80	8.55	64.48	-1.39	27.47	6.25	67.9560
4742	Istanbul	Turkey	72.75	38.79	31.58	35.59	2.63	186.67	16.51	28.25	16.90	53.0525
4521	Sydney	Australia	18.25	186.45	27.36	45.47	8.91	170.28	-0.68	90.85	5.72	71.0560

Next steps: [Generate code with df_reduced](#) [New interactive sheet](#)

4) Data Scaling & Discretization

Before:

	City	Country	PM2.5	PM10	NO2	SO2	CO	O3	Temperature	Humidity	Wind Speed	AQI
6252	Beijing	China	111.92	32.08	7.95	37.45	9.43	45.77	23.68	30.31	5.91	55.3610
4684	Berlin	Germany	62.62	153.88	22.76	22.12	0.45	112.53	1.93	66.61	1.22	71.6845
1731	Bangkok	Thailand	77.06	114.31	24.99	45.80	8.55	64.48	-1.39	27.47	6.25	67.9560
4742	Istanbul	Turkey	72.75	38.79	31.58	35.59	2.63	186.67	16.51	28.25	16.90	53.0525
4521	Sydney	Australia	18.25	186.45	27.36	45.47	8.91	170.28	-0.68	90.85	5.72	71.0560

After:

```
[37] # 4) Data Scaling and Discretization
✓ 0s
scaler = MinMaxScaler()
num_cols = ["PM2.5", "PM10", "NO2", "SO2", "CO", "O3", "AQI"]
df_reduced[num_cols] = scaler.fit_transform(df_reduced[num_cols])

# Discretize AQI into categories
df_reduced["AQI_Category"] = pd.cut(
    df_reduced["AQI"],
    bins=[0, 0.3, 0.6, 0.8, 1.0],
    labels=["Good", "Moderate", "Unhealthy", "Hazardous"]
)
df_reduced.head()



|      | City     | Country   | PM2.5    | PM10     | NO2      | SO2      | CO       | O3       | Temperature | Humidity | Wind Speed | AQI      | AQI_Category |
|------|----------|-----------|----------|----------|----------|----------|----------|----------|-------------|----------|------------|----------|--------------|
| 6252 | Beijing  | China     | 0.737390 | 0.115983 | 0.030983 | 0.744181 | 0.942424 | 0.187997 | 23.68       | 30.31    | 5.91       | 0.369557 | Moderate     |
| 4684 | Berlin   | Germany   | 0.397226 | 0.757812 | 0.187059 | 0.431196 | 0.035354 | 0.539754 | 1.93        | 66.61    | 1.22       | 0.512216 | Moderate     |
| 1731 | Bangkok  | Thailand  | 0.496861 | 0.549297 | 0.210560 | 0.914659 | 0.853535 | 0.286580 | -1.39       | 27.47    | 6.25       | 0.479630 | Moderate     |
| 4742 | Istanbul | Turkey    | 0.467122 | 0.151341 | 0.280008 | 0.706207 | 0.255556 | 0.930397 | 16.51       | 28.25    | 16.90      | 0.349382 | Moderate     |
| 4521 | Sydney   | Australia | 0.091078 | 0.929441 | 0.235536 | 0.907922 | 0.889899 | 0.844038 | -0.68       | 90.85    | 5.72       | 0.506723 | Moderate     |


```

4. Statistical Analysis

1) General overview

```
▶ # General overview
df_reduced.info()



| Index: 2000 entries, 6252 to 6929 |              |                |          |
|-----------------------------------|--------------|----------------|----------|
| Data columns (total 13 columns):  |              |                |          |
| #                                 | Column       | Non-Null Count | Dtype    |
| 0                                 | City         | 2000 non-null  | object   |
| 1                                 | Country      | 2000 non-null  | object   |
| 2                                 | PM2.5        | 2000 non-null  | float64  |
| 3                                 | PM10         | 2000 non-null  | float64  |
| 4                                 | NO2          | 2000 non-null  | float64  |
| 5                                 | SO2          | 2000 non-null  | float64  |
| 6                                 | CO           | 2000 non-null  | float64  |
| 7                                 | O3           | 2000 non-null  | float64  |
| 8                                 | Temperature  | 2000 non-null  | float64  |
| 9                                 | Humidity     | 2000 non-null  | float64  |
| 10                                | Wind Speed   | 2000 non-null  | float64  |
| 11                                | AQI          | 2000 non-null  | float64  |
| 12                                | AQI_Category | 1999 non-null  | category |



dtypes: category(1), float64(10), object(2)  
memory usage: 205.3+ KB


```

[40] df_reduced.describe()

	PM2.5	PM10	NO2	S02	CO	O3	Temperature	Humidity	Wind Speed	AQI
count	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	0.501342	0.503600	0.491887	0.494081	0.507387	0.508792	14.43329	55.030790	10.205715	0.495096
std	0.288942	0.293132	0.285086	0.287523	0.286817	0.291054	14.50690	26.102797	5.633121	0.180474
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-9.97000	10.010000	0.500000	0.000000
25%	0.254675	0.253557	0.245732	0.239026	0.263384	0.246733	1.42750	32.827500	5.237500	0.366880
50%	0.501069	0.503452	0.481768	0.488975	0.508081	0.518125	14.43500	54.990000	10.290000	0.497188
75%	0.746447	0.756626	0.736379	0.751837	0.760859	0.762264	27.11250	77.675000	14.900000	0.623202
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	39.97000	99.970000	19.990000	1.000000

2) Central Tendency Measures

[44] # Central Tendency Measures

```
mean_pm25 = df_reduced["PM2.5"].mean()
median_pm25 = df_reduced["PM2.5"].median()
mode_pm25 = df_reduced["PM2.5"].mode()[0]
min_pm25 = df_reduced["PM2.5"].min()
max_pm25 = df_reduced["PM2.5"].max()

print("PM2.5 - \n Min:", min_pm25,
      "\n Max:", max_pm25,
      "\n Mean:", mean_pm25,
      "\n Median:", median_pm25,
      "\n Mode:", mode_pm25)
```

PM2.5 -

Min: 0.0
Max: 1.0
Mean: 0.5013420961843649
Median: 0.5010694818188091
Mode: 0.9833022838611745

3) Dispersion Measures

```
[47] ✓ 0s ➔ # Dispersion Measures
      range_pm25 = max_pm25 - min_pm25
      variance_pm25 = df_reduced["PM2.5"].var()
      std_pm25 = df_reduced["PM2.5"].std()
      Q1_pm25, Q3_pm25 = df_reduced["PM2.5"].quantile([0.25, 0.75])
      IQR_pm25 = Q3_pm25 - Q1_pm25

      print(" Range:", range_pm25,
            "\n Variance:", variance_pm25,
            "\n StdDev:", std_pm25,
            "\n IQR:", IQR_pm25)

 ➔ Range: 1.0
 Variance: 0.08348766571104879
 StdDev: 0.2889423224642745
 IQR: 0.4917718898778721
```

4) Correction Analysis

